

STM-based concurrent heaps

Mateusz Machalica

Heap (a priority queue)


- for simplicity: min-heap
- operations on heap of size n
 - `min()` in $O(1)$ time
 - `push(e)` in $O(\log(n))$ time
 - `pop()` in $O(\log(n))$ time, possibly amortized

Take 1: coarse-grained locking

- take array-based heap implementation
- replace array with `TArray Int e`
 - equivalently with `MArray Int (TVar e)`
- replace `size` with `TVar Int`
- put everything in a transaction

Take 1: correctness (trivial invariant)

“A heap has the heap property after each transaction.”




$$p \leq \min(l, r)$$

Take 1: problems

- full contention on `size` required for `push()` and `pop()`
- consecutive `push()` operations follow nearly the same path
 - despite inserted values
 - long operations shuffle many elements
 - STM creates data dependencies for elements on the path
 - two updates cannot “follow” each other

Take 1: concurrent push ()

- STM has no idea of our invariant
- altered elements will force second transaction to be rejected
- full sequencing



Problems to keep in mind

- STM thinks that if element is touched then it necessarily breaks the transaction
 - very strong invariant
 - but not always true
 - and hard to weaken
- arrays
 - cannot make things immutable
 - but a tree would be nicer for dynamic sized heaps
- updating `size` cannot end up in a global transaction

Assumption

“Mutual exclusion between `push()` and `pop()` operations.”

In case we drop it:

- what is the semantics?
- how to avoid contention in the root?
- hard to maintain any invariants
- not necessary for most applications I can think about...



Take 2: fine-grained locking

- implement heap as a tree
 - all operations go top-down
 - balancing achieved by maintaining `size` of a subtree rooted at each node
- single transaction covers $O(1)$ nodes
 - we need some fancy invariant
 - if the old one works, we can sort in $O(n)$ time in comparison model `#impossibru`

Why it *should* work

- updates can proceed one after another on the same path
 - we do not look back after the transaction
 - each update waits for preceding to leave the root
 - k updates take $O(k)$ time if $\log(n) = O(k)$ and $k = O(\log(n))$
- or diverge as early as possible
 - each update descends in a different subtree
 - k updates take $O(k)$ time if $\log(n) = O(k)$ and $k = O(n)$
- in sequential case
 - k updates take $O(k \log(n))$

Why it *should* work



The push () operation

- in transaction
 - choose a child node with smaller `size` field
 - update `size`
 - if the node is empty, insert our value and finish
 - swap element to be inserted e with the one in current node x if $e < x$
- recurse into the chosen node
 - after committing the transaction

Correctness of `push()`


- invariant of a transaction
 - the simple invariant works for the tree, but need to cover the extra element e
 - “The heap property holds for each tree node and the element to be inserted e is greater than all elements on the path from the current node to the root”
- in absence of `pop()` operations
 - elements in each node are non-increasing
 - traversing a node = it will be at most as big as we have seen until the next deletion

The pop () operation

- in transaction
 - replace root's element with “a gap”
 - decrease `size`
- in transaction
 - choose a child node with smaller element
 - decrease `size`
 - move element in the node to the parent, place gap here
- recurse into the chosen node
 - after committing the transaction
 - only if it's subtree is non-empty

The pop () operation - caveat

- if a root has a gap `retry`
- if any of the child nodes have gaps and `size > 0`, then `retry`
 - there is at least one node having gap-free children
 - no deadlock
- are these necessary?
 - we don't know which element will end up here
 - possibly the smallest among children of the node



Correctness of pop ()


- invariant of a transaction
 - the simple invariant no longer works!
 - “The heap property holds for each tree node, when one assigns each gap-node an element from it’s parent (recursively)”
- for each gap we have a processor trying to push it down the tree
- the conditions from previous slide ensure
 - that we know upper bound on each subtree before replacing root’s gap
 - we never steal the gap from someone else

Performance comparison


- initially 50k elements in the heap
- each (of W) writer inserts $50k / W$
- each (of R) reader removes $50k / R$
- different scenarios (parametrized by N)
 - $W = \max(1, N - 1), R = 1$
 - $W = N, R = 0$
 - $W = 0, R = N$
 - $W = 1, R = 1$, sequentially, i.e. first all writes, then all removals
- averaged over 5 random seeds

Performance comparison

CoarseHeap performance



FineHeap performance



Application: concurrent sorting

- insert all elements into a heap
 - using N processors
 - each inserts around n/N elements
- initialize n element TArray Int (Maybe e) with Nothing
- assign threads to array cells in round-robin fashion
 - waits for previous cell to be filled
 - replaces heap's root with a gap, places the element in the array
 - pushes the gap down the tree
 - proceeds to it's next index

Performance: concurrent sorting

- 100k random elements
- phase 1 (insertion)
 - each (of N) writer inserts $100k / N$ elements
- phase 2 (removal)
 - each (of N) reader
 - waits for previous cell of output array to be written
 - takes element from the heap
 - writes to the output array into its cell
 - cells assigned to workers in round-robin fashion

Performance: concurrent sorting

